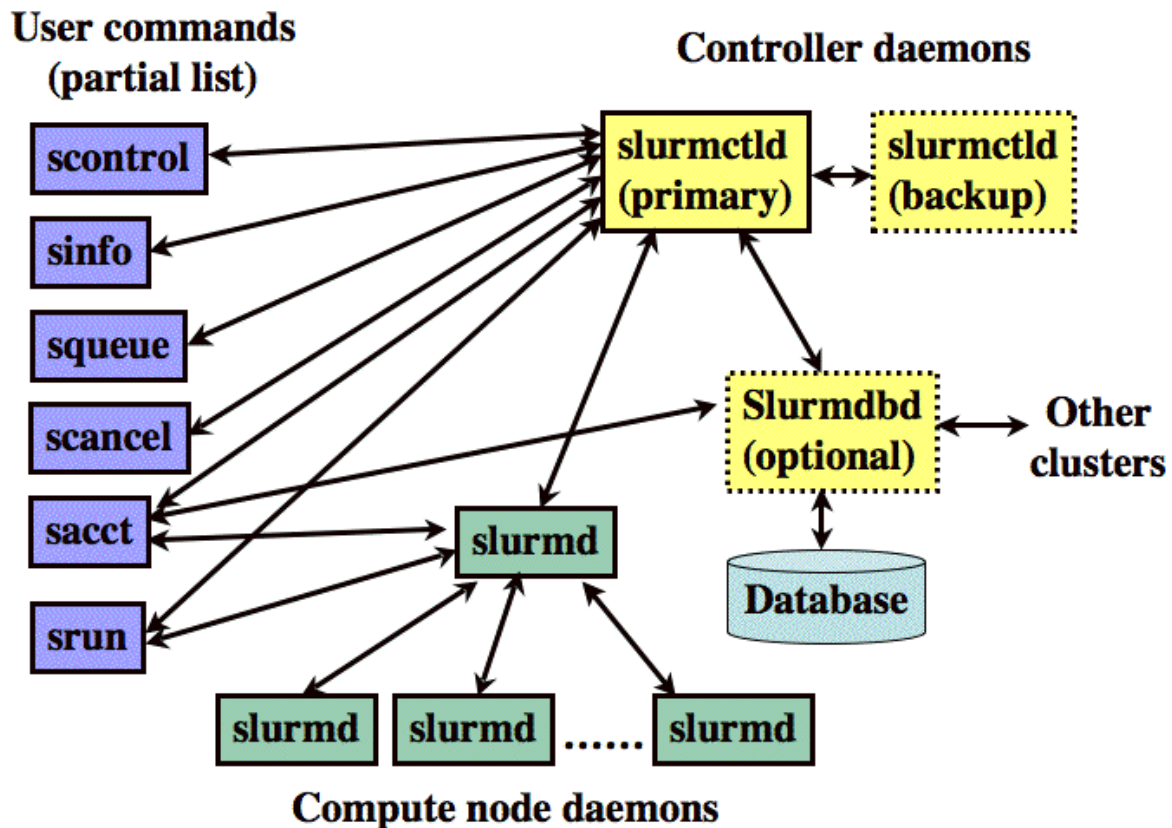


Slurm Installation



Slurmctld-slurm controller
Node heat
Scheduling
Node demon communication
DBD

SYSTEM REQUIREMENT

create 3 vm of centos image

Server Type	Server Hostname	Specs	
Master	slurm_master	2GB Ram,	2vcpus
Worker	slurm_node1	2GB Ram,	1vcpus
Worker	slurm_node2	2GB Ram,	1vcpus

edit host file and write all ip-address of all machine on all system

vi /etc/hosts

```
127.0.0.1    localhost localhost.localdomain localhost4 localhost4.localdomain4
::1         localhost localhost.localdomain localhost6 localhost6.localdomain6
192.168.174.166 master
192.168.174.167 node1
192.168.174.168 node2
```

make sure all 3 pc ping to each other

create keygen on master node

Ssh-keygen

copy keygen to all client node

Ssh-copy-id

Update all system

yum update

disable selinux on all system

setenforce 0

'or'

vi /etc/selinux/config

SELINUX=disabled

```
# This file controls the state of SELinux on the system.
# SELINUX= can take one of these three values:
#   enforcing - SELinux security policy is enforced.
#   permissive - SELinux prints warnings instead of enforcing.
#   disabled - No SELinux policy is loaded.
SELINUX=disabled
# SELINUXTYPE= can take one of three values:
#   targeted - Targeted processes are protected,
#   minimum - Modification of targeted policy. Only selected processes are protected.
#   mls - Multi Level Security protection.
SELINUXTYPE=targeted
```

firewall disable on all system

```
systemctl disable firewalld
systemctl stop firewalld
```

restat all system on all machine

```
init 6
```

```
on master:-
```

```
-----
```

```
yum install nfs-utils           #install nfs-utils
yum install nfs4* -y           #install nfs-utils version 4
rpm -qa | grep nfs             #show nfs file which u
install
mkdir shital                   #create a directory
cd shital
    touch file1 file2          #create file in directory
systemctl start nfs             #start nfs service
systemctl stop firewalld        #stop firewall service
vi /etc/exports                 #create a file
    /root/shital *(rw,sync,no_root_squash) #write this in file as same
systemctl restart nfs           #restart nfs service
exportfs -ar                    #show file export or not
```

```
on client:-
```

```
-----
```

```
yum install nfs-utils           #install nfs-utils
yum install nfs4*               #install nfs-utils version 4
rpm -qa | grep nfs              #show nfs file which u
install
mkdir surya
cd shital
    touch file3 file4
systemctl start nfs
systemctl stop firewalld
mount -t nfs master:/root/shital /root/surya #mount client and master
directory
```

```
ls
file1 file2 file3 file4           #show on both system
```

Install chrony(for set time)

```
-----
```

On server—

install chrony

yum install chrony* -y

vi /etc/chrony.conf

26(write server-ip network) & 29 → uncomment

Ex-192.168.174.0/24

local stratum 10

```
# Use public servers from the pool.ntp.org project.
# Please consider joining the pool (http://www.pool.ntp.org/join.html).
server 192.168.100.129 iburst
#server 1.centos.pool.ntp.org iburst
#server 2.centos.pool.ntp.org iburst
#server 3.centos.pool.ntp.org iburst

# Record the rate at which the system clock gains/losses time.
driftfile /var/lib/chrony/drift

# Allow the system clock to be stepped in the first three updates
# if its offset is larger than 1 second.
makestep 1.0 3

# Enable kernel synchronization of the real-time clock (RTC).
rtcsync

# Enable hardware timestamping on all interfaces that support it.
#hwtimestamp *

# Increase the minimum number of selectable sources required to adjust
# the system clock.
#minsources 2

# Allow NTP client access from local network.
allow 192.168.100.0/24

# Serve time even if not synchronized to a time source.
local stratum 10

# Specify file containing keys for NTP authentication.
#keyfile /etc/chrony.keys

# Specify directory for log files.
logdir /var/log/chrony

# Select which information is logged.
#log measurements statistics tracking
~
~
"/etc/chrony.conf" 38L, 1105C
```

:wq

copy the file to all node

```
[root@master]# rsync /etc/chrony.conf node1:/etc/chrony.conf
[root@master]# rsync /etc/chrony.conf node2:/etc/chrony.conf
```

start chrony service

systemctl start chronyd

On client:-

install chrony

yum install chrony* -y

start chrony service

systemctl start chronyd

for request to server time query

ntpdate -q 'server-ip' -

for request to server time update

ntpdate -u 'server_ip' -

check all machine time status

chronyc sources

```
[root@node1 ~]# ntpdate -q 192.168.174.174
server 192.168.174.174, stratum 10, offset 0.540681, delay 0.02629
16 Jan 18:35:00 ntpdate[10759]: step time server 192.168.174.174 offset 0.540681 sec
[root@node1 ~]# ntpdate -u 192.168.174.174
16 Jan 18:35:29 ntpdate[10760]: step time server 192.168.174.174 offset 0.540677 sec
[root@node1 ~]# chronyc sources
210 Number of sources = 4
MS Name/IP address          Stratum Poll Reach LastRx Last sample
=====
^? 64.227.167.110            0 10    0    -    +0ns[ +0ns] +/- 0ns
^? 162.159.200.123           0 10    0    -    +0ns[ +0ns] +/- 0ns
^? 178.215.228.24            0 10    0    -    +0ns[ +0ns] +/- 0ns
^? 192.46.215.60             0 10    0    -    +0ns[ +0ns] +/- 0ns
[root@node1 ~]#
```

download slurm tar file from browser

<https://download.schedmd.com/slurm/slurm-22.05.7.tar.bz2>

install mariaDB on all machine

```
yum install mariadb-server mariadb-devel -y
```

epel-release on all machine

```
yum install epel-release -y
```

install munge on all machine

```
yum install munge munge-libs munge-devel -y
```

rpm-build on all machine

```
yum install rpm-build -y
```

rpm build of slurm file on master

```
rpmbuild -ta slurm-22.05.7.tar.bz2
```

install perl-extutils on master

```
yum install perl-ExtUtils-Install* -y
```

```
yum install python3 readline-devel pam-devel gcc -y
```

Generate munge key on master

```
/usr/sbin/create-munge-key -r
```

change permission of munge.key on master

```
chmod 400 /etc/munge/munge.key
```

send key to all node

```
scp munge.key root@node1:/etc/munge/
```

```
scp munge.key root@node2:/etc/munge/
```

change owner on all node

```
chown -R munge: /etc/munge/munge.key
```

start and enable munge service

```
systemctl enable munge  
systemctl start munge  
systemctl status munge
```

Go to this path

```
cd /root/rpmbuild/RPMS//x86-64
```

Install slurm on all node

```
yum install slurm* -y
```

copy all rpm file to all node

```
scp -r /root/rpmbuild/RPMS/x86_64/ root@node1:/home/  
scp -r /root/rpmbuild/RPMS/x86_64/ root@node2:/home/
```

Export slumuser on all node

```
export SLURMUSER=992
```

```
groupadd -g $SLURMUSER slurm
```

```
useradd -m -c "SLURM workload manager" -d /var/lib/slurm -u  
$SLURMUSER -g slurm -s /bin/bash slurm
```

change file name on master

```
cp /etc/slurm/slurm.conf.example /etc/slurm/slurm.conf
```

edit file on master (line no 11&12&52&57)

```
vi /etc/slurm/slurm.conf
```

```
clusterName=surya                #controlmachine=surya  
SlurmctldHost=master             #controlAddr=master  
SlurmdSpoolDir=/var/share/slurm  
StateSaveLocation=/var/share/slurm/ctld
```


create directory on master

```
mkdir -p /var/share/slurm/ctld
```

create file on master

```
touch /var/log/slurmctld.log
```

Give permission to ctld file

```
chown -R slurm:slurm /var/share/slurm
```

on all client

```
slurmd -C
```

```
[root@node1 suryal]# slurmd -C
NodeName=node1 CPUs=1 Boards=1 SocketsPerBoard=1 CoresPerSocket=1 ThreadsPerCore=1 RealMemory=1959
UpTime=0-02:40:58
```

copy 1st line

On master:-

```
vi /etc/slurm/slurm.conf
```

```
# COMPUTE NODES
#NodeName=linux[1-32] CPUs=1 State=UNKNOWN
NodeName=node2 CPUs=1 Boards=1 SocketsPerBoard=1 CoresPerSocket=1 ThreadsPerCore=1 RealMemory=1959
NodeName=node1 CPUs=1 Boards=1 SocketsPerBoard=1 CoresPerSocket=1 ThreadsPerCore=1 RealMemory=1959
PartitionName=debug Nodes=ALL Default=YES MaxTime=INFINITE State=UP
```

hint:check 2nd last line

copy slurm.conf file on all node

```
scp /etc/slurm/slurm.conf root@node1:/etc/slurm/slurm.conf
```

```
scp /etc/slurm/slurm.conf root@node2:/etc/slurm/slurm.conf
```

```
[root@master ~]# scp /etc/slurm/slurm.conf root@node1:/etc/slurm/slurm.conf
slurm.conf                                100% 3229      2.4MB/s   00:00
[root@master ~]# scp /etc/slurm/slurm.conf root@node2:/etc/slurm/slurm.conf
slurm.conf                                100% 3229      2.6MB/s   00:00
```

change owner of slurm.conf on all node

```
chown -R slurm:slurm /etc/slurm/slurm.conf
```

On client:-

make directory

```
mkdir -p /var/share/slurm/d
```

change owner

```
chown -R slurm:slurm /var/share/slurm/
```

create file

```
touch /var/log/slurmd.log
```

start munge service

```
systemctl start munge
```

start slurmd

```
systemctl start slurmd
```

enable slurmd

```
systemctl enable slurmd
```

run on server for check their worker nodes active or not

```
sinfo
```

```
[root@master ~]# sinfo
PARTITION AVAIL  TIMELIMIT  NODES  STATE NODELIST
debug*     up       infinite    1  idle*  node1
debug*     up       infinite    1  down*  node2
```

for start forcefully node

```
scontrol update nodename='down_node_name' state=idle
```

provide shell of the node

```
srun -N1 -pty /bin/sh
```

show node on server

scontrol show node 'node_name'

show partition on server

scontrol show part 'part_name'

-----X-----X-----X-----X-----

Create partition on server

vi /etc/slurm/slurm.conf

```
PartitionName=debug Nodes=ALL Default=YES
MaxTime=INFINITE State=UP
PartitionName=small Nodes=node1 Default=YES
MaxTime=INFINITE State=UP
PartitionName=medium Nodes=node2 Default=YES
MaxTime=INFINITE State=UP
```

```
PartitionName=debug Nodes=ALL Default=YES MaxTime=INFINITE State=UP
PartitionName=small Nodes=node1 Default=YES MaxTime=INFINITE State=UP
PartitionName=medium Nodes=node2 Default=YES MaxTime=INFINITE State=UP
```

add line like this in last row

copy slurm.conf file on both nodes

```
[root@master ~]# scp /etc/slurm/slurm.conf node1:/etc/slurm/slurm.conf
slurm.conf                                100% 3371      3.0MB/s   00:00
[root@master ~]# scp /etc/slurm/slurm.conf node2:/etc/slurm/slurm.conf
Warning: Permanently added the ECDSA host key for IP address '192.168.174.180' to the list
of known hosts.
slurm.conf                                100% 3371      2.2MB/s   00:00
[root@master ~]# systemctl restart slurmctld
[root@master ~]# systemctl status slurmctld
```

restart slurmctld on master

systemctl restart slurmctld

systemctl status slurmctld

```
# restart slurmd on nodes
systemctl restart slurmd
systemctl status slurmd
```

run on server for check their worker nodes status
sinfo

```
[root@master ~]# sinfo
PARTITION AVAIL  TIMELIMIT  NODES  STATE NODELIST
debug      up       infinite    2    idle node[1-2]
small      up       infinite    1    idle node1
medium*    up       infinite    1    idle node2
```

Idle- ready to start for run any job

Up- node start

Down-node down

Comp-submitted their job